Kafka Streaming in Java

Objective: To send desired data from a CSV file via Java as a Producer for Kafka, which will be consumed by Java Consumer

STEPS:

Step 1: Start Zookeeper

Command: ‘zkserver’

Step 2: To Start Kafka Server

Command: “ .\bin\windows\kafka-server-start.bat .\config\server.properties”

Note: Change your directory while entering command to Kafka folder

Step 3: Creating a Topic for Kafka to which messages will be send and retrieve

Command: “kafka-topics.bat --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic producerTest”

“producerTest” is the name of our Topic

Note: Change your directory to /bin/windows while entering the command

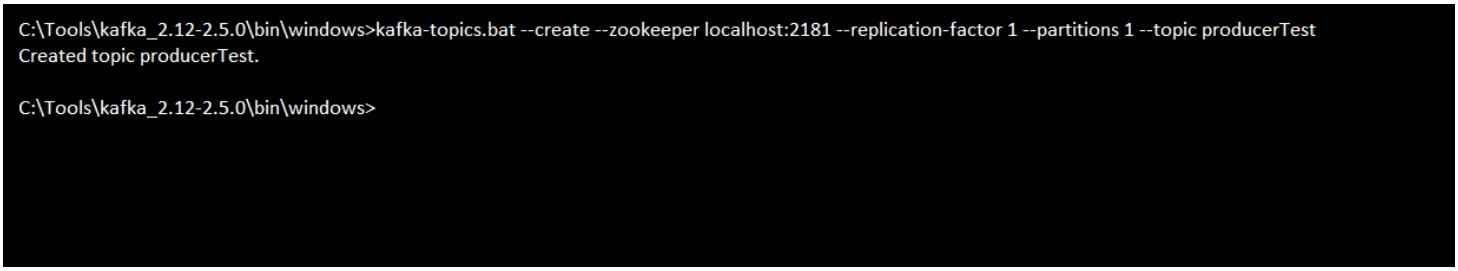

The default port number for zookeeper: 2181

Replication factor: 1

Partition made: 1

Step 4: Creating a Java Producer

Package Manager: Since we are using Java we will be using Maven as our package manager to download dependencies, using Intellij Idea

Adding Java Kafka Client dependencies to .iml/.xml file. These dependencies are to be added in order to run the program

```
<?xml version="1.0" encoding="UTF-8"?>
<module type="JAVA_MODULE" version="4">
 <component name="NewModuleRootManager" inherit-compiler-output="true">
  <exclude-output />
  <content url="file://$MODULE_DIR$">
   <sourceFolder url="file://$MODULE_DIR$/src" isTestSource="false" />
  </content>
  <orderEntry type="inheritedJdk" />
  <orderEntry type="sourceFolder" forTests="false" />
  <orderEntry type="library" name="org.apache.kafka:kafka_2.12:0.10.1.1" level="project" />
  <orderEntry type="library" name="org.apache.kafka:connect-api:1.0.1" level="project" />
  <orderEntry type="library" name="com.opencsv:opencsv:5.1" level="project" />
 </component>
</module>
```

Configuration needed to connect to Kafka Servers:

- We need some properties that will help us to connect to Kafka server.
- These properties will be key-value pair.

bootstrap.servers:

- This property will define the address of the Kafka server.
- It is the combination of IP address and Port on which Kafka server is hosted.
- Default value: localhost:9092.

key.serializer:

- Kafka stores the data as key-value pair in Topics.
- We need to define the type of serializer that will be needed to convert the content of key so that it could be sent to Kafka server.
- If key is of String type, then we need to set its value as org.apache.kafka.common.serialization.StringSerializer

value.serializer:

- We need to define the type of serializer that will be needed to convert the content of value so that it could be sent to Kafka server.
- If value is of String type, then we need to set its value as org.apache.kafka.common.serialization.StringSerializer

```
Properties properties = new Properties();
properties.put("bootstrap.servers", "localhost:9092");
properties.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
properties.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
```

Publish data to Kafka server:

- ProducerRecord type object is created which is send to the Kafka Server

```
ProducerRecord producerRecord = new ProducerRecord(topic_name, key, value)
```

- Topic name is where data will be published, data can be consumed from topic.
- Key inside which data is stored.
- Value is the actual data.

Note: We will avoid making a key in our program

Connection to Kafka and sending message:

- We need to make a connection to the Kafka server using KafkaProducer class.
- KafkaProducer takes Properties object as an argument to create a connection.

```
KafkaProducer kafkaProducer = new KafkaProducer(properties);
```

- We can send data to Kafka server by using send command.

```
kafkaProducer.send(new ProducerRecord("producerTest",finalData));
```

- producerTest is the name of Topic.
- finalData is the value of data to be published.

```
kafkaProducer.close();
```

- To close the connection, we created.

CSV:

Producer will publish the data to Kafka server by reading it from a CSV file:

```
String path = " kafka_data\\team_info.csv";
FileReader filereader = new FileReader(path);
CSVReader csvReader = new CSVReader(filereader);
String[] nextRecord;
```

- path is a String type variable containing the address of a CSV file.
- Opening CSV file by creating an object of FileReader class.
- Passing the FileReader object as an argument while creating CSVReader object.
- csvReader is the object containing the data of a CSV file.
- nextRecord is a String array which will contain one line of data at a time.

```
int counter = 0;
System.out.println("Enter the line number from which you want to publish data: ");
int lineNumStartFrom = sc.nextInt();
System.out.println("\nEnter the line number to which you want to publish data: ");
int lineNumTo = sc.nextInt();
while ((nextRecord = csvReader.readNext()) != null) {
   if (counter > lineNumStartFrom && counter <= lineNumTo) {
      String data = Arrays.toString(nextRecord);
      String finalData = data.substring(1, data.length()-1);
      kafkaProducer.send(new ProducerRecord("producerTest",finalData));
      Thread.sleep(1000);
   }
   counter++;
}
```

Final Code:

```
import java.util.Arrays;
import java.util.Scanner;
import java.io.FileReader;
import java.util.Properties;
import com.opencsv.CSVReader;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;

public class Producer {
   public static void main(String[] args) {
      Scanner sc = new Scanner(System.in);
      Properties properties = new Properties();
      properties.put("bootstrap.servers", "localhost:9092");
      properties.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
      properties.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
      KafkaProducer kafkaProducer = new KafkaProducer(properties);
      try {
         String path = " kafka_data\\team_info.csv";
         FileReader filereader = new FileReader(path);
         CSVReader csvReader = new CSVReader(filereader);
         String[] nextRecord;
         int counter = 0;
         System.out.println("Enter the line number from which you want to publish data: ");
         int lineNumStartFrom = sc.nextInt();
         System.out.println("\nEnter the line number to which you want to publish data: ");
         int lineNumTo = sc.nextInt();
         while ((nextRecord = csvReader.readNext()) != null) {
            if (counter > lineNumStartFrom && counter <= lineNumTo) {
               String data = Arrays.toString(nextRecord);
               String finalData = data.substring(1, data.length()-1);
               kafkaProducer.send(new ProducerRecord("producerTest",finalData));
               Thread.sleep(1000);
            }
            counter++;
         }
      }
      catch (Exception e) {
         e.printStackTrace();
      }
      finally {
         kafkaProducer.close();
      }
   }
}
```

Step 5: Creating a Java Consumer

Package Manager: Since we are using Java we will be using Maven as our package manager to download dependencies, using Intellij Idea

Adding Java Kafka Client dependencies to .iml/.xml file. These dependencies are to be added in order to run the program

```
<?xml version="1.0" encoding="UTF-8"?>
<module type="JAVA_MODULE" version="4">
 <component name="NewModuleRootManager" inherit-compiler-output="true">
  <exclude-output />
  <content url="file://$MODULE_DIR$">
   <sourceFolder url="file://$MODULE_DIR$/src" isTestSource="false" />
  </content>
  <orderEntry type="inheritedJdk" />
  <orderEntry type="sourceFolder" forTests="false" />
  <orderEntry type="library" name="org.apache.kafka:connect-api:1.0.1" level="project" />
 </component>
</module>
```

Connection Details:

- topic variable contains the name of Topic which we are subscribed to.
- bootstrap.servers, key.deserializer, value.deserializer are explained in producer.
- ‘enable.auto.commit’, ‘true’ means consumer will commit automatically on receiving messages.
- auto.commit.interval.ms = 1,000ms
- session.timeout.ms = 30,000ms

```
String topic = "producerTest";
String group = "producerTest";
Properties properties = new Properties();
properties.put("bootstrap.servers", "localhost:9092");
properties.put("group.id", group);
properties.put("enable.auto.commit", "true");
properties.put("auto.commit.interval.ms", "1000");
properties.put("session.timeout.ms", "30000");
properties.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
properties.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
```

Connection to Kafka and sending message:

- We need to make a connection to the Kafka server using KafkaConsumer class.
- KafkaConsumer takes Properties object as an argument to create a connection.

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<String, String>(properties);
```

Subscribing to the Topic

```
consumer.subscribe(Arrays.asList(topic));
```

Printing messages received from selected Topic by subscribing it

```
while (true) {
   ConsumerRecords<String, String> records = consumer.poll(100);
   for (ConsumerRecord<String, String> record : records)
      System.out.printf(record.value()+"\n");
}
```

Final Code:

```
import java.util.*;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;

public class Consumer {
   public static void main(String[] args) throws Exception {

      String topic = "producerTest";
      String group = "producerTest";
      Properties properties = new Properties();
      properties.put("bootstrap.servers", "localhost:9092");
      properties.put("group.id", group);
      properties.put("enable.auto.commit", "true");
      properties.put("auto.commit.interval.ms", "1000");
      properties.put("session.timeout.ms", "30000");
      properties.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
      properties.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
      KafkaConsumer<String, String> consumer = new KafkaConsumer<String, String>(properties);

      consumer.subscribe(Arrays.asList(topic));
      System.out.println("Subscribed to topic " + topic);

      while (true) {
         ConsumerRecords<String, String> records = consumer.poll(100);
         for (ConsumerRecord<String, String> record : records)
            System.out.printf(record.value()+"\n");
      }
   }
}
```

Final Output:

Java Output-IntelliJ (Producer):

user can decide the amount of data he needs to publish

```
Enter the line number from which you want to publish data:
5

Enter the line number to which you want to publish data:
15

Process finished with exit code 0
```

Java Output-IntelliJ (Consumer):

published data will be updated in every 1000ms

```
Subscribed to topic producerTest
5, 17, Pittsburgh, Penguins, PIT, /api/v1/teams/5
17, 12, Detroit, Red Wings, DET, /api/v1/teams/17
28, 29, San Jose, Sharks, SJS, /api/v1/teams/28
18, 34, Nashville, Predators, NSH, /api/v1/teams/18
23, 20, Vancouver, Canucks, VAN, /api/v1/teams/23
16, 11, Chicago, Blackhawks, CHI, /api/v1/teams/16
9, 30, Ottawa, Senators, OTT, /api/v1/teams/9
8, 1, Montreal, Canadiens, MTL, /api/v1/teams/8
30, 37, Minnesota, Wild, MIN, /api/v1/teams/30
15, 24, Washington, Capitals, WSH, /api/v1/teams/15
```